

Builders Pitch Fest 2026

Greyin

The AI-native trust and verification layer for hiring experienced talent

6 core features | 3 revenue lines | One shared Hub identity | <https://greyin.net>

Startup Snapshot

1. What problem are you trying to solve

Senior, experienced professionals (12+ years) get lost in over-crowded hiring platforms, filtered out by keyword-matching ATS software, and made redundant by restructuring -- not because they stopped being good at the work, but because nothing proves it at scale. Greyin is a single-identity ecosystem built for senior professionals, where an AI-driven verification and matching layer turns real work -- a completed project, a real review, a real deliverable -- into a portable, audited trust score employers can search against, for hiring and adjacent talent decisions.

2. What inspired you to work on this problem?

The gap became obvious building each core feature as its own product first: a hiring platform where “verified” meant nothing beyond a resume upload, a freelance marketplace with no real proof of delivery. Manual review doesn't scale past a handful of profiles -- AI is what makes verifying thousands of real interactions, and matching against them accurately, actually possible.

Problem Understanding

1. What problem are you solving, and who experiences it most acutely?

Professionals with 12-30+ years of experience feel it most acutely: overqualified for junior ATS filters, with no structured way to prove a real track record short of a reference call. Employers feel the inverse -- unable to tell a genuinely Verified Expert from a well-formatted resume until well into an expensive interview process, and with no visibility into who's a fit for a role that doesn't exist yet.

2. What research, observations, or evidence validate this opportunity?

Built directly from this platform's own construction: every feature that launched with a self-reported trust signal (an unverified "connection," a self-picked seniority flag) had to be re-audited and re-gated to a real, AI-and-human-verified interaction once real usage began. That re-audit is the direct evidence: unverified social proof is the default failure mode of every adjacent product, and only a verification layer with AI doing the heavy lifting can close it at any real scale.

Customer & Opportunity

1. Who is your ideal customer, and who would make the buying decision?

The paying customer is the corporate employer: a hiring manager, talent-acquisition lead, or Head of Talent at a company that regularly hires senior/fractional talent and is tired of paying recruiter fees or burning interview cycles on unverifiable resumes. Individual professionals (12+ years -- senior ICs, directors, fractional executives) are the free, identity-led supply side; the enterprise buyer is who Greyin is built to monetize, budget-line by budget-line: recruiting spend, workforce-planning spend, and L&D/mentorship spend.

2. How large is the opportunity if your assumptions prove correct?

India alone has tens of millions of professionals past the 12-year mark and a multi-billion-dollar enterprise recruiting spend category already. Three separate employer budget lines are addressable: recruiting spend redirected from traditional job boards and recruiter fees toward a **Verified Expert** search subscription; workforce-planning spend for Longlist's anticipatory pipeline, a category with no real incumbent today; and freelance/advisory spend redirected from 20%+-commission marketplaces toward FlexPro's low-fee model -- a flat subscription plus a modest service fee split across both sides, not a 20%+ cut of every transaction.

Proposed Solution

1. What is your proposed solution, and how does it address the identified problem?

Greyin is one platform with 6 core features, not 7 separate products -- Greyin Hub is the shared identity layer that binds them, not a feature in itself. Of the 6, 3 are revenue lines billed straight to the corporate employer (one, Longlist, stands alone; two live inside existing features). The other 3 produce the quality profiles and verified interactions that feed the Greyin Score -- the actual signal being sold. See the breakdown below.

Category	Name	Hosted In	Role
Revenue line	Outplacement	DeepEdge	Enterprise pays per-headcount -- departing employees routed into the Verified Expert ecosystem
Revenue line	Coaching (Mentorship)	FlexPro, surfaced platform-wide	Verified Experts monetize expertise directly -- sessions, cohorts, packages; platform take-rate
Feature -- Primary	Longlist (Future Hiring Pipeline)	own app	Enterprise pays for AI-matched, anticipatory access 3-12 months before a role opens
Feature -- Supporting	DeepEdge (core hiring)	own app	Verified job board + candidate database -- every real hire strengthens the Score
Feature -- Supporting	FlexPro	own app	Subscription + low-fee gig delivery -- feeds verified-outcome reviews into the Score
Feature -- Supporting	StackWorks	own app	Real project delivery, AI+human verified -- the single strongest Score input
Feature -- Supporting	Salt & Pepper + GreyMatters	own apps	Community engagement + published expertise -- reputation signal, organic reach

2. What are the key capabilities you envision?

A single AI-verified identity across all 6 core features via Greyin Hub; an AI-and-human-verified interaction gate on every social-proof signal; a Bayesian-shrunk Greyin Score counting only real work; and AI matching that turns feature activity directly into revenue-line accuracy.

3. What measurable value do you expect the solution to deliver?

Three distinct paid budget lines for employers; one compounding score for individuals; a free Supporting tier whose real output -- score integrity -- is what the Primary tier actually monetizes. Every score is explainable to the person it's about -- per-pillar inputs, evidence counts and weights are visible on their profile.

Technology & AI Approach

1. What role will AI play in your proposed solution, and why is it essential?

The mechanism that makes “verified” scale past manual review -- what employers actually pay for. Six AI systems run in production today, plus four on the near-term roadmap (marked below):

2. What technologies, AI models, or frameworks do you plan to use, and why?

A pluggable, admin-configurable LLM registry (Anthropic, OpenAI, self-hosted Ollama fallback) behind every feature -- never locked to one vendor; every call is timeout-bounded and fails over automatically instead of breaking a paying customer's workflow.

AI Capability	What It Does	Status
AI-Verified Search	NL query → structured, ranked candidate filters	LIVE
Verified-Outcome Scoring	AI+human review of real StackWorks deliverables	LIVE
Longlist AI Matching	Surfaces future-role fits, no application needed	LIVE
Quality & Moderation Gates	Pre-publish safety gate (blocks before insert) + periodic quality sweep, across 3 features	LIVE
Gig-Listing Quality Assist	On-demand AI rewrite tips at post time -- informational, never gates publish	LIVE
Multi-Provider Core	Anthropic / OpenAI / Ollama, no vendor lock-in	LIVE
Explainable Match Rationale	Plain-language “why this match” per result	ROADMAP
Bias / Fairness Auditing	Active skew monitoring on AI match output	ROADMAP
AI Mentor Matching	Best-fit coach suggestions -- Score + goals	ROADMAP
AI Reference Synthesis	Multi-response consistency reports for employers	ROADMAP

3. What makes you feel your product is a genuine AI product, and not an AI wrapper?

Every call changes what data exists afterward, feeding the next match -- not a chat layer wrapped around someone else's directory. Each verified interaction compounds into every future match's accuracy: a real data network effect.

Competitive Landscape

1. What alternatives currently exist?

LinkedIn and Naukri/Indeed for hiring; Upwork and Toptal for freelance work; Blind and niche communities for peer discussion; ADPList and informal referral networks for mentorship. Each solves one slice in isolation, and every one of them treats “trust” as a self-reported field an AI can't meaningfully verify because there's no real interaction underneath it to check.

2. Why do you believe your approach is better suited to solving this problem?

Three things a generic AI wrapper can't touch. The test: remove the AI -- does the product still work?

(1) Compounding data: the graph underneath our AI is real and gated against self-report -- a competitor can copy a search UI in a quarter, not manufacture years of verified interaction history.

(2) A new category: Longlist matches people against roles that don't exist publicly yet -- delete the AI and this feature is not degraded, it's impossible; there is no manual version of anticipatory matching at scale.

(3) Trustable output: every signal feeding the AI requires a real interaction gate, unlike endorsements an AI can rank but never verify -- fewer bad hires from AI-assisted screening, a compliance argument for the enterprise buyer, not a UX one.

	Typical Social Network	Greyin
Core loop	Connections, attention, engagement	Verified work -> AI-scored trust -> AI-matched revenue
What the AI does	Feed ranking, engagement optimization	Verification scoring + anticipatory matching
What's monetized	Attention/ads, access to a network	AI-verified matches and AI-synthesized trust
Remove the AI	Network still works -- AI is decorative	Longlist and AI-Verified Search cannot function

Business Model & Go-To-Market

1. Who do you expect will pay for your solution?

The corporate employer pays across three Primary tiers, each a distinct budget line, not one subscription trying to cover everything.

(1) Outplacement: sales-led, invoiced per headcount -- already live intake + admin triage (enterprise_leads pipeline), tapping restructuring/HR-transition budget.

(2) Coaching: Verified Experts sell mentor sessions, cohorts, packages and recordings directly -- already live with real Razorpay checkout -- tapping L&D budget, Greyin takes a platform commission.

(3) Future Hiring Pipeline (Longlist): AI-matched, anticipatory access to talent 3-12 months out -- an Enterprise-tier add-on to DeepEdge's existing live subscription for regular hiring (server-side price-validated) -- tapping workforce-planning budget, a category with no real incumbent today.

AI capability is the explicit upsell lever across all three.

Individual membership stays free: the Supporting tier only compounds the Score these three sell if supply-side adoption is unrestricted.

2. How do you plan to acquire your first customers?

Direct enterprise outreach for Outplacement and Coaching (HR/L&D buyers, sales-led by design); Longlist as the wedge into a full Enterprise contract -- a company testing the AI-matched pipeline for one future role is a warm path to all three tiers; and Supporting-tier community growth (Salt & Pepper, GreyMatters, the built-in +25-score referral program) as the free acquisition engine that keeps the Score -- and therefore all three Primary tiers -- credible.

Development Roadmap

1. What are the next major milestones toward building your MVP?

All 6 core features (plus Greyin Hub, the shared identity layer) are already built and live, pre-launch on purpose while the trust/security layer was hardened first (38-finding audit, fully closed). The monetization path from here is a phased pilot-to-scale roadmap, not one launch date:

2. What resources or support will be most critical during this journey?

Design-partner employers, enterprise-sales expertise, and infrastructure support scaled only once real load justifies it.

Phase	Timeframe	Focus & Expert Supply	Revenue Milestone (Target)
Pre-Launch Hardening	Mo. -6 to -3	Security/trust-layer audit (done, 38 findings closed); competitive-parity pass (done -- employer self-service, community events, feedback + wishlist loops, unified brand across all 6 apps, cold-start seed content); design-partner outreach; invite-only expert cohort	\$0 by design — no live transactions yet
Closed Pilot	Mo. -3 to 0	3-5 design-partner employers on DeepEdge + Longlist; curated 50-100 expert cohort; validate AI-match quality	First LOIs signed; pricing validated with pilots
Public Launch	Mo. 0-1	Open signup, all 6 core features; referral-driven expert acquisition (+25 score/conversion); GreyMatters content marketing	First paying cohort converts from pilot
Year 1	Mo. 1-12	Scale employer accounts; AI-Verified Search + Longlist AI-matching activated broadly; mentor commission live	Target: first ~50 paying employer accounts
Year 2	Year 2	Enterprise-tier accounts for high-volume teams; StackWorks verified-outcome premium tooling; geo expansion	Target: 3-5x employer base; first Enterprise contracts
Year 3	Year 3	Vertical expansion; reference-check API for enterprise HR systems; verified-interaction data moat compounds	Target: sustainable unit economics; Series A readiness

Team

1. Why is your team uniquely positioned to solve this problem?

Solo founder. 13+ years in automotive embedded systems -- most of it in verification and validation: professionally proving that safety-critical software does what its spec claims, and scaling V&V teams 3-4x for Stellantis followed by 2+ years of freelancing after a pivot into Full-stack/AI Engineering.

Greyin turns that same discipline on professional trust -- don't take the claim, verify the interaction underneath it. I'm also the user: after a deliberate pivot into AI and two years shipping full-stack and AI software solo, I re-entered the hiring market as a senior candidate and hit the exact wall this platform removes -- a resume upload is the only proof a decade of real work gets, and keyword filters do the rest. One of those solo builds was itself a hiring portal; watching "verified" mean nothing on it is what turned six separate products into one platform built around a verification layer.

2. What relevant domain or technical expertise does the team possess?

Full-stack ownership of the entire platform as built: schema design and row-level-security-first data architecture, 7 independently-deployable Next.js applications, a pluggable multi-provider LLM/AI integration layer powering 6 live AI systems, and a self-run production deployment -- plus a genuinely rigorous engineering practice: a real 38-finding security audit closed and individually verified end to end, and a feature-by-feature teardown against a parallel AI-generated build of the same concept.

Why BITSoM Vertex?

1. Why have you applied to the BITSoM Vertex programme?

BITSoM Vertex's combination of Silicon Valley AI venture-building expertise, enterprise pilot pathways, and investor access is exactly the missing piece: the AI infrastructure and the product are built and live, but converting that into a funded, enterprise-validated company needs precisely the venture-building rigor and warm enterprise introductions this programme provides.

2. What is the biggest challenge preventing you from building your MVP?

Not technical execution -- the platform, its 6 live AI systems, and its security posture are already real. The gap is deliberately commercial: converting a solid, pre-launch AI product into signed design-partner pilots and a first paying enterprise cohort, exactly the enterprise-pathway support this programme offers.

Supporting Material

Provide links to your:

- Prototype (if available)

Live production platform -- all 6 core features plus Greyin Hub deployed and reachable today.

- Demo Video (if available)

Not yet recorded -- available on request ahead of the Mumbai Pitch Fest.

- GitHub Repository (if available)

<https://github.com/zollyzeus/greyin> (private repo); access + architecture walkthrough on request.

- Website / Landing Page (if available)

<https://greyin.net> (ecosystem hub / landing page)

<https://deepedge.greyin.net> | <https://flexpro.greyin.net> | <https://longlist.greyin.net>

<https://saltnpapper.greyin.net> | <https://stackworks.greyin.net> | <https://greymatters.greyin.net>

- Contact Details

Anand Rajendran, Founder

+91 98843 99915 | rajendran.anand@live.com | <https://www.linkedin.com/in/anand4585> | <https://anandrajendran.dev>

Builders Pitch Fest 2026

THANK YOU